Milestone: 20GB Reached

We've reached 20GB of stored data — thank you for helping us grow!

Patreon Ko-fi Liberapay
Close
Freedium

- ko-fi
- librepay
- patreon

< Go to the original

Meta's AI-Scored Engineers: "Use AI or You're Out"

If you're shipping more than 50% of your code by hand without AI, you're out.

Agent Native Follow
~7 min read · February 4, 2026 (Updated: February 4, 2026) · Free: No

> *If you're shipping more than 50% of your code by hand without AI, you're out.*

Meta is tying p**erformance and bonuses more tightly to AI-driven impact**, and if you think this stays at Meta, you haven't been paying attention.

Amazon has already required employees to show how they use **AI in promotion packets** in certain units.

Cisco published internal research suggesting employees who use AI more often are more **likely to be promoted and more likely to be labeled critical to retain.**

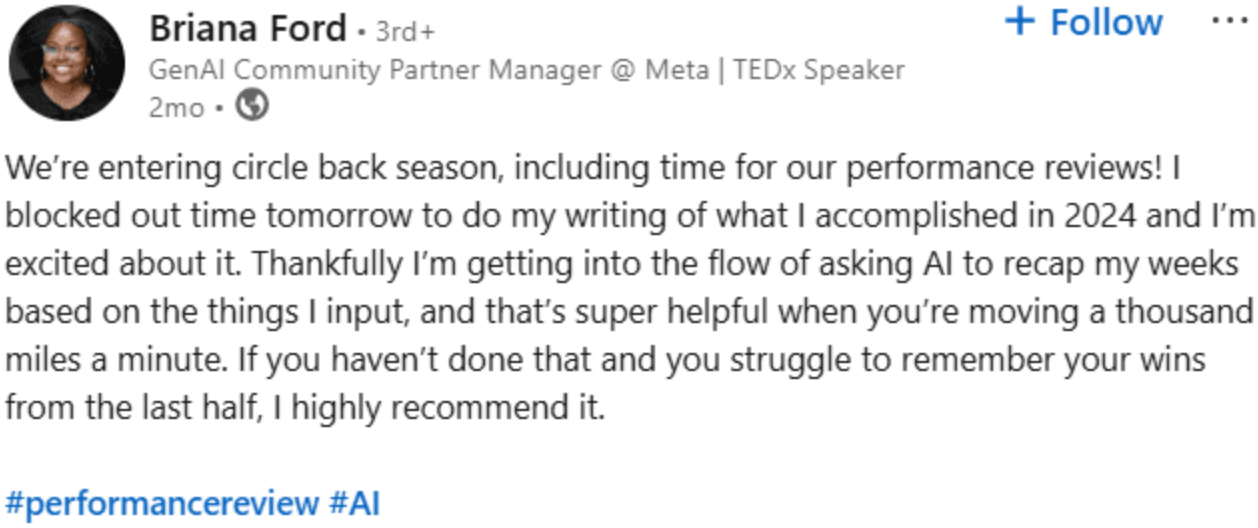

Source: Linkedin

> If your company is moving in this direction, drop a comment with what's changing, i.e. how it's measured, what it rewards, and what it breaks.

So what should you do, **as an employee, as a developer, as an engineering or product leader**, when the performance game itself is changing?

Because engineering productivity has been measured for years.

We've had DORA metrics, we've had SPACE, we've had endless debates about what "good" even means: **lead time vs reliability**, **throughput vs burnout**, **shipping vs quality**.

Those frameworks are imperfect, but at their best they try to measure **outcomes at the team/system level** and they come with context, caveats, and humility about what numbers can't capture.

This wave is different.

Now introduce a **bonus multiplier**. Now introduce **forced ranking**. Now introduce the promise that **AI-driven impact** gets rewarded.

People have always gamed fragile systems and that's where the fairness problem explodes, who did they think was going to win?

And meanwhile, many of you have already wired agents into everything: repo access, docs, calendar, inbox, **the full "I'm so back" stack**.

So here are the only two questions that matter:

- What happens when the way engineering labor is evaluated changes?
- What do systems do when you attach real money to proxies?

Why this is spreading now

If you read this article today, please think about the **economics and org design** of your company.

Mark Zuckerberg expects 2026 to be a year where AI "**dramatically change[s] the way that we work**," describing investment in "AI-native tooling," "flattening teams," and **projects that used to require big teams now being done by a single very talented person**.

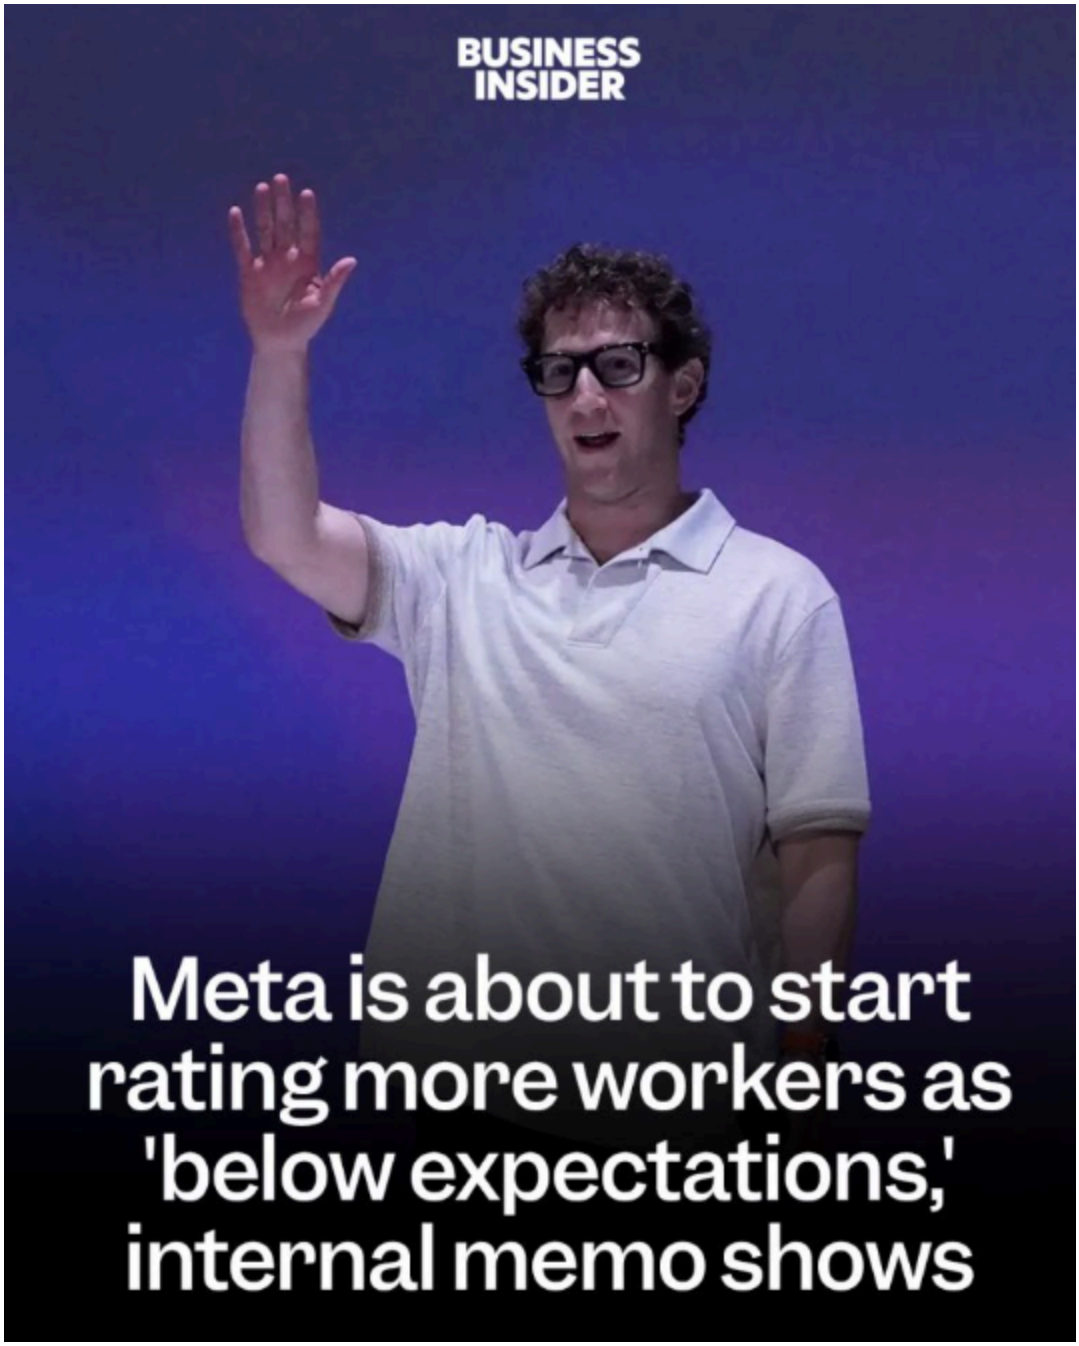

Read that again.

Flattening teams + elevating individual contributors + AI-native tooling is the new operating model.

And that model needs a new management system because when teams get flatter and individuals carry larger scopes, managers want "evidence" that output is real, consistent, and attributable.

So you instrument. You summarize. You score.

The temptation is **evaluation of AI productivity**.

Your job becomes legible to the tracker

Most engineering impact is not immediately legible to dashboards.

The best work often looks like:

- deleting code
- preventing incidents
- simplifying architectures
- reducing dependencies
- making on-call survivable
- building primitives other teams reuse
- tightening security posture
- designing interfaces that eliminate classes of bugs

None of that naturally shows up as "more generated lines."

So when a system starts sampling your work from repositories, docs, and activity trails, and compressing it into a summary, it will bias toward what's easiest to capture.

That creates a new kind of pressure:

> Don't just do valuable work.
>
> Do valuable work in a way that produces machine-readable evidence.

This will change behavior.

The "data points" question is the trap

So you may ask, "What are the data points?".

Gilad Naor • 3rd+
Ex-Meta. Land a FAANG role or the promotion that you deserve.
2mo •

+ Follow ···

At Meta, the performance review process works something like this:

-Everyone writes a self assessment and sends it to their manager.
-Every manager write an assessment packet for each of their directs. This includes a proposed rating.
-Managers meet with managers of similar teams to compare notes and ensure that their proposed rating is fair. It is very common for the proposed ratings to change at this step.
-All the managers of a large org meet together to review all the packets and ensure that the proposed ratings makes sense. It is very common for the proposed ratings to change at this step.
-There are follow-ups on flagged edge cases that require more investigation.
-After some broad sanity checks by senior leadership and HR, the ratings are locked.

Meta reportedly tracks 200+ data points in this system.

Do we know the exact list? No.

But if you've built internal analytics systems, you can guess the categories:

- PRs opened/merged
- review comments given/received
- cycle time
- diff size
- bug/incident associations
- test changes
- doc edits
- meeting artifacts
- task/ticket completions
- AI tool usage frequency
- "AI generated vs hand-authored" proportions

Issue here is that the more data points you collect, the more you create the illusion of objectivity while quietly increasing the surface area for misinterpretation.

The model understands correlates, and **correlates are gameable**.

Which brings us to the new meta-game.

The new elite won't be best engineers

If you tie compensation to AI-driven impact, three tribes will appear:

1) Builders who use AI the way we all want this to work: accelerate drafts, explore options, reduce toil, tighten feedback loops, ship better systems.

2) Performers who optimize for optics: tokens, artifacts, visible diffs, frequent commits, lots of generated content, lots of activity evidence.

3) **Invisible** who does risk reduction, reliability, mentoring, architecture review, security fixes, migrations, and operational hardening.

This is ruthless but also clarifying, because it exposes what many orgs are actually doing, they're redefining baseline competence as AI-mediated production.

Where productivity claims go to die

Speeding up generation doesn't speed up delivery if verification doesn't scale.

AI can increase the volume of candidate code, candidate docs, candidate plans, but shipping still requires:

- review
- tests
- security checks
- architecture consistency
- operational readiness
- rollout discipline
- monitoring and incident response

So the bottleneck moves.

From writing to reviewing, from coding to auditing, from producing to verifying.

If leadership only measures the **front of the pipeline**, it will misread what's happening.

And teams will feel it as a tax:

- more regressions
- more flaky tests
- more review fatigue
- more subtle security mistakes
- more "looks right" code that's wrong in production

The security is where it gets ugly

Now we layer in the second half of the original context: agentic tools.

A bunch of people joked about installing OpenClaw/Clawdbot/Moltbot-like **agents as productivity engines, burning tokens, getting promoted twice**, etc.

Funny.

This is exactly how you end up with shadow IT except now it's shadow agents, and shadow agents are worse, because they can act.

So here's the paradox companies are creating:

1. "Use agents to get more done."
2. "We will track whether you use them."
3. "Agents expand attack surface and can be exploited."
4. "We have not fully solved safety, sandboxing, and governance."
5. "Your bonus depends on adoption anyway."

Incentives win.

People will use the tools, route around friction, and if official tools are locked down, they'll use unofficial ones.

This is how you get:

- prompt injection incidents
- credential leaks
- "agent did something weird" outages
- internal data exposure
- supply-chain style compromises via plugins/skills

And it will happen faster than most governance teams can respond, because performance incentives can move behavior in weeks.

A single "AI usage" rubric is incompatible with diverse high-impact work

Research, infra, security, reliability, privacy, architecture, and product strategy do not all express impact through the same observable signals.

If you score everyone by one proxy, you get conformity, and conformity is how innovation dies

But there is one way this can work and it's not what most companies are doing.

If you want AI-heavy development to be real, you need an operational stack like:

- **Provenance:** require authors to declare where AI was used and why
- **Policy-as-code:** linting, secret scanning, dependency checks, license compliance
- **Test gates:** minimum coverage expectations for risky modules, property-based tests where it fits
- **Diff constraints:** limit AI-generated diff size unless a human provides rationale
- **Rollout discipline:** canary + monitoring + fast rollback as default
- **Agent sandboxing:** least privilege tokens, scoped access, strong audit logs, blocked dangerous actions by default
- **Compute budgets:** token spend as a real cost center, not "free magic"

Notice what's missing?

A bonus tied to "how many AI-generated lines you produced."

So what should leaders measure instead?

You can measure outcomes which are hard to game and aligned with value:

- incident reduction
- MTTR improvement
- customer-impacting bug rate
- lead time improvements without rollback spikes
- operational toil reduction
- security posture improvements (least privilege, audit coverage, policy compliance)
- adoption of reusable primitives by other teams

and use input metrics only as diagnostics and never for compensation:

- tool usage
- token consumption
- AI vs human authored proportions
- number of agent runs
- "AI-generated LOC"

You can almost treat these like CPU metrics, e.g. useful for debugging, disastrous as a performance score.

What you should do if your company moves this way?

You don't need to become a token-maximizer.

You need to become a person who turns AI into **durable outcomes** that are legible:

- Keep a running **impact log** (incidents prevented, risk reduced, systems simplified)
- **Quantify** what matters (time-to-detect, time-to-recover, support load, cost)
- Make your **reasoning visible** (design docs, tradeoffs, constraints)
- Use AI as a **multiplier** but build the verification muscle (tests, reviews, threat modeling)
- Avoid **usage theater**, because the rules will change once leadership realizes it's being gamed

And if your organization starts valuing usage over results, consider what that says about the long-term culture.

The best people often leave first in environments where metrics replace judgment.

Bonus Articles

Local LLMs That Can Replace Claude Code

Small team of engineers can easily burn >$2K/mo on Anthropic's Claude Code (Sonnet/Opus 4.5). As budgets are tight, you...

medium.com

"This PDF hates OCR" - DeepSeek OCR 2 Reads Documents Like A Human

"This PDF hates OCR" - DeepSeek OCR 2 Reads Documents Like A Human Your customer forwards an invoice PDF with the...

medium.com

Vectorless RAG for Agents: PageIndex Is Why Their Demo Works and Yours Needs "Context"

When I'm building a RAG apps, the thing that quietly kills quality (and budgets) is turning a document into a pile of...

medium.com

Fully Local Agentic Coding on localhost: Claude Code, Codex, llama.cpp, Ollama and Docker MR

This is the setup I wish I had the first time I tried to run Claude Code and Codex-style workflows against local...

medium.com

Open Source Perplexity: Perplexica is Self-Hosted AI Search With Citations

Perplexica is an open-source, AI-powered search engine inspired by Perplexity-style UX.

medium.com

Qwen3-Max-Thinking Outperforms Claude Opus 4.5 and Gpt-5.2: The Panic Is Real

Alibaba’s Qwen3 family of models has become a major alternative in the LLM ecosystem.

medium.com
#ai-productivity#developer-productivity#performance-reviews#ai-coding#coding

Reporting a Problem

Sometimes we have problems displaying some Medium posts.

If you have a problem that some images aren't loading - try using VPN. Probably you have problem with access to Medium CDN (or fucking Cloudflare's bot detection algorithms are blocking you).

Problem Description Describe your problem here...

Submit Cancel