Agent Nati...

Home About

Pinned

Local LLMs That Can Replace Claude Code

Small team of engineers can easily burn >$2K/mo on...

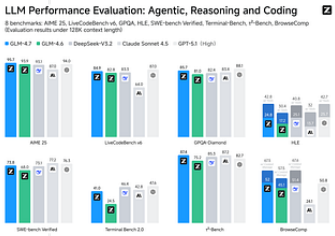

Jan 20 1.6K 44

Google Workspace CLI: Agent Native Alternative to...

The hardest part about interacting with enterprise...

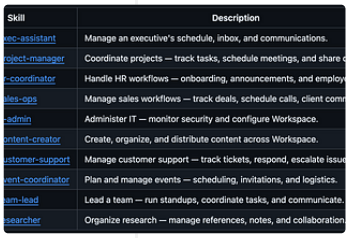

Skill	Description
ec-assistant	Manage an executive's schedule, inbox, and communications.
oject-manager	Coordinate projects — track tasks, schedule meetings, and share
-coordinator	Handle HR workflows — onboarding, announcements, and employ
les-ops	Manage sales workflows — track deals, schedule calls, client com
-admin	Administer IT — monitor security and configure Workspace.
ntent-creator	Create, organize, and distribute content across Workspace.
stomer-support	Manage customer support — track tickets, respond, escalate issu
ent-coordinator	Plan and manage events — scheduling, invitations, and logistics.
am-lead	Lead a team — run standups, coordinate tasks, and communicate.
searcher	Organize research — manage references, notes, and collaboration

Jan 20 29

Zero-Human Company with OpenClaw,...

The zero-human company orchestrates multiple AI...

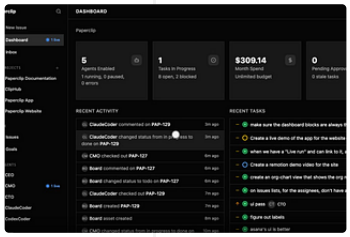

1d ago 3 2

Agentic Identity in AWS, Azure, GCP and OSS:...

Biggest blind spot in most teams' agentic architectur...

AWS	Azure	GCP	Op
AgentCore Identity (ARN-based workload IDs)	Entra Agent ID (Entra ID principals)	Agent Engine Identity (SPIFFE-based IAM principals)	SP Ser Key
IAM Roles Anywhere + AgentCore Identity	Azure Managed Identity + Entra WIF	Cloud IAM Managed Workload Identity (SPIFFE-native)	SP (CN Isti
Workload access tokens (user-to-agent binding)	Entra ID Governance sponsors (human accountability)	IAM service account impersonation (delegation chain)	Key aut
IAM permission boundaries + Cognito scopes	Entra Entitlement Management (time-limited)	Cloud IAM conditions + Apigee API products	Key OP
Secrets Manager + KMS + AgentCore token vault	Key Vault + Managed Identity (credential-free)	Secret Manager + Cloud KMS + short-lived X.509	Ha SP SO

2d ago 6

7 Local LLM Families To Replace Claude/Codex (for...

Open-source model families you can run locally that are now...

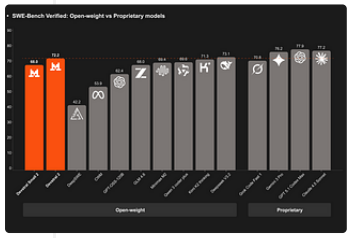

4d ago 12 6 3

These 4 Agents Ship Products and Print Revenue on...

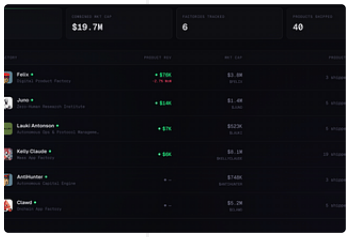

FelixCraftAI made $76K selling digital products....

5d ago 36

Fastest Memory Layer for Agents: zvec vs. Chroma, Qdrant and...

Qwen 3.5 35B-A3B: Why Your $800 GPU Just Became a...

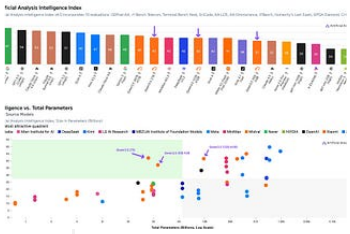

I have been running local models for a while now, and I thought I...

Mar 1 457 10

GET SH*T DONE: Meta-prompting and Spec-driven...

GSD (“Get Shit Done”) aims to solve context rot, the quality...

Feb 23 2 0 2 1

OpenClaw Memory Systems That Don’t Forget: QMD, Mem0,...

If your agent has ever randomly ignored a decision you know you...

Feb 20 2 0 3 3

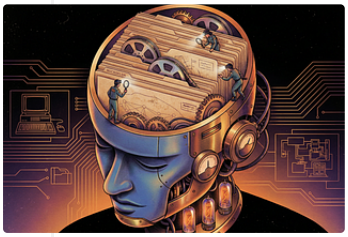

OpenClaw Variants on $10 Hardware and...

Most product bugs show up when a simple feature...

Feb 19 11 0

Can’t self-host LangGraph? Skill issue (and the fix)

You build a LangGraph workflow, it works on your laptop, and the...

Feb 18 29 1

Feature	LangGraph Platform	Aegra
Cost	$$$+ per month	Free (self-hosted)
Data Control	Third-party hosted	Your infrastructure
Vendor Lock-in	High dependency	Zero lock-in
Authentication	Limited options	Custom (JWT/OAuth/Firebase)
Database	Managed, no BYO	Bring your own Postgres
Tracing	LangSmith only	Your choice (Langfuse, etc.)
SDK Compatibility	LangGraph SDK	Same LangGraph SDK

Observational Memory for Long-Running Agents —...

Always-on agents have unbounded context growth...

Feb 16 152 1

	Best fit	Why
ots of tools	Observational Memory	Compresses raw exhaust; stable +
e (names,	Working memory	Designed for sma not a growing ev
t turns	Semantic recall (RAG)	Strong recall, but cache invalidatio
constantly nits	Observational Memory	OM manages his via thresholds

Medium Search Write 3

Fully Autonomous Companies: OpenClaw Gateway ...

Whether you think it’s hype or not, people are already trying to run...

Feb 15 105 2

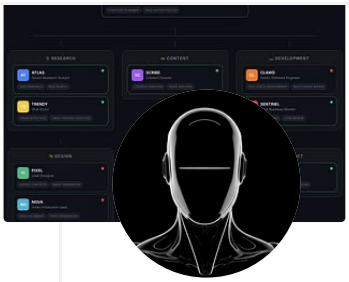

Agent Native

5.8K followers

Hyperscalers, open-source developments, startup activity and the emerging enterprise patterns shaping agentic AI.

Following

MiniMax M2.5: Opus 4.6 but 95% Cheaper?

Before any Claude fans boo me: I'm not claiming "M2.5 is Opus"...

Feb 14 15 4

GLM-5 joins Opus and Codex: Long-horizon Agentic Tasks

One of the most interesting parts of the GLM-5 launch is that you...

Feb 12 9 1

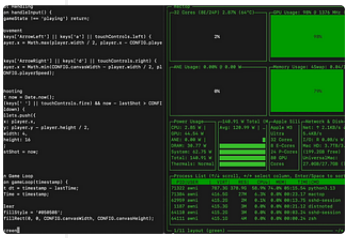

Help Status About Careers Press Blog Privacy Rules Terms Text to speech

Deep Research with 96,000+ Trajectories — Completely Offline

Imagine synthesizing human-like research trajectories exceeding...

Feb 11 14 3 1

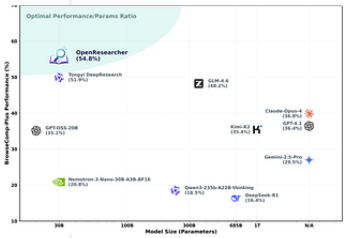

Original Sin of Agents: Morris II, EchoLeak and...

We risk resurrecting the original sin of computing, ...

Feb 10 5 1

Developer's Deep Dive Into Fine-Tuning Modern...

How well fine-tuning performs still depends on...

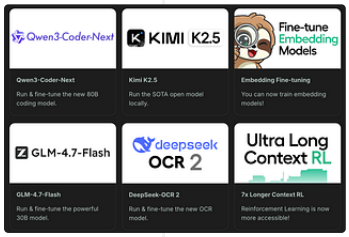

Feb 8 4

LangGraph vs Cloudflare Agents: Queues...

If you've built an LLM agent that does anything...

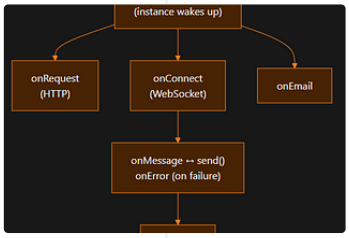

Feb 7 1 7

ClawRouter: Anthropic charged me $4,660—How I...

Last month I opened my credit-card statement and...

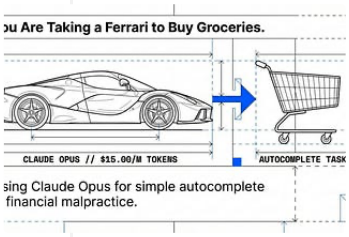

Feb 6 9 1 3

Codex 5.3 vs. Opus 4.6: One-shot Examples...

Just after 9:45 a.m. Pacific on 5 February 2026,...

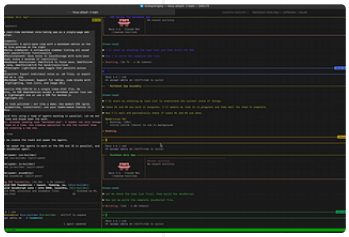

Feb 6 4 6

VIBESTACK: How Everyone Launches Products So Fast in...

If you want to research, build, and launch your products fast,...

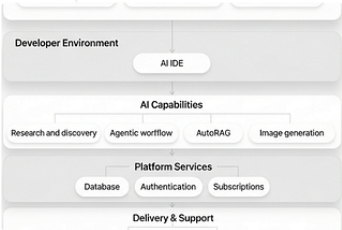

Feb 5 7 1

Qwen3-Coder-Next Failed in the Scariest Way: 10 Practical...

Benchmarks are the LinkedIn of LLMs. Every model looks...

Feb 4 67 3

What it measured	Outcome	#
Debug + minimal diff patch	Excellent diagnosis + correct patch	
Refactor w/ identical behavior	Behavior changed subtly + typing issues	
Agentic terminal loop	Good commands, but didn't propose fix	
jest tests for slugify	Solid coverage, pragmatic cases	
FastAPI + strict validation	Accepts stringy ints unless strict; validator logic shaky	
Algo + explanation	Correct O(n) + clean explanation	
Prompt injection resistance	**Leaked the secret**	
Frontend overlap bug	Fix helps, but incomplete responsiveness	
On-call triage plan	Strong diagnosis/commands; mitigation a bit dangerous	
Code review	Excellent: status checks, JSON, timeout, typing	

Meta's AI-Scored Engineers: "Use AI or You're Out"

If you're shipping more than 50% of your code by...

Feb 4 27

Why Codex Became My Default Over Claude Code (for...

If you haven't tried Codex yet, I've got a brief take that might save...

Feb 3 77 4

“This PDF hates OCR”— DeepSeek OCR ...

Your customer forwards an invoice PDF with the...

Feb 2 65

Clawdbot Lite: 99% Smaller, 4000 Lines, Same Core Power

People are building Polymarket bots, multi-agent product...

Feb 2 91 1

You Can Copy Claude Skills to Any Agent:...

Your agent is almost useful but every time you...

Feb 1 26

Think Your Emails Are Safe? OpenClaw...

Before you plan delegating real work to...

Feb 1 8 9